

# CSCI5502\_Project\_Sharma\_Tapas\_Xia

April 23, 2024

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from imblearn.over_sampling import SMOTE
from google.colab import drive
drive.mount('/content/drive')
data_path = r'./drive/My Drive/creditcard.csv'
df = pd.read_csv(data_path)
```

Mounted at /content/drive

```
[ ]: creditcard_df = pd.read_csv(data_path)
```

```
[ ]: unprocessed_df = creditcard_df
```

**#Dataset Description Details:-** The dataset comprises credit card transactions conducted by European cardholders in September 2013. This dataset comprises transactions that took place during a span of two days, with a total of 492 instances of fraud out of a total of 284,807 transactions.

The dataset comprises numerical input variables that have undergone a Principal Component Analysis (PCA) transformation. Some of the original characteristics are `distance_from_home`, `distance_from_last_transaction`, `ratio_to_median_purchase_price`, `repeat_retailer`, `used_chip`, `used_pin_number`, and `online_order`, which are transformed into variables using PCA. The principle components derived with PCA are denoted as V1, V2,... V28. The only characteristics that have not undergone PCA transformation are 'Time' and 'Amount'. Characteristic In the dataset, the variable 'Time' represents the duration in seconds between each transaction and the initial transaction. The 'Amount' feature represents the transaction amount and can be utilized for example-dependent cost-sensitive learning. Characteristic The response variable, denoted as 'Class', assumes a value of 1 when fraud is present and 0 when it is not.

The dataset exhibits a significant imbalance, with the positive class (defined as frauds) representing a mere 0.172% of the total transactions. There are total 284,807 records and 31 fields.

To solve this Imbalanced issue we will implement **SMOTE** algorithm to make the transaction balanced. Dataset contains numerical input variables which are the result of a PCA transformation. In the original dataset we

**Source** - <https://data.world/raghu543/credit-card-fraud-data>

**Solving method:-** The given problem statement is comes under binary classification We have to solve problem using different machine learning algorithm as well as deep learning algorithms

**Limitations** - Due to confidentiality issues, the initial characteristics and additional contextual details of the data are converted into major components features V1, V2,... V28 using PCA.

```
[ ]: creditcard_df.describe()
```

```
[ ]:
```

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.759061e-12	-8.251130e-13	-9.654937e-13	8.321385e-13
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	1.649999e-13	4.248366e-13	-3.054600e-13	8.777971e-14	-1.179749e-12
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	V21	V22	V23	V24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	-3.405756e-13	-5.723197e-13	-9.725856e-13	1.464150e-12
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	-6.987102e-13	-5.617874e-13	3.332082e-12	-3.518874e-12	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

```

                class
count  284807.000000
mean      0.001727
std       0.041527
min       0.000000
25%      0.000000
50%      0.000000
75%      0.000000
max       1.000000

```

```
[8 rows x 31 columns]
```

```
[ ]: creditcard_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    284807 non-null   float64
1   V1       284807 non-null   float64
2   V2       284807 non-null   float64
3   V3       284807 non-null   float64
4   V4       284807 non-null   float64
5   V5       284807 non-null   float64
6   V6       284807 non-null   float64
7   V7       284807 non-null   float64
8   V8       284807 non-null   float64
9   V9       284807 non-null   float64
10  V10      284807 non-null   float64
11  V11      284807 non-null   float64
12  V12      284807 non-null   float64
13  V13      284807 non-null   float64
14  V14      284807 non-null   float64
15  V15      284807 non-null   float64
16  V16      284807 non-null   float64
17  V17      284807 non-null   float64
18  V18      284807 non-null   float64
19  V19      284807 non-null   float64
20  V20      284807 non-null   float64
21  V21      284807 non-null   float64
22  V22      284807 non-null   float64
23  V23      284807 non-null   float64
24  V24      284807 non-null   float64
25  V25      284807 non-null   float64
26  V26      284807 non-null   float64
27  V27      284807 non-null   float64

```

```
28 V28      284807 non-null float64
29 Amount  284807 non-null float64
30 class   284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

All columns have either float or integer datatypes, indicating the absence of noise values such as ? or .. If such values were present, the datatype would have been **object**.

Removing all the null values from the dataset and replacing those values with mean value of the column.

```
[ ]: for column in creditcard_df.columns:
      if creditcard_df[column].isna().sum() > 0:
          mean_value = creditcard_df[column].mean()
          creditcard_df[column].fillna(mean_value, inplace=True)
```

```
[ ]: creditcard_df.isna().sum()
```

```
[ ]: Time      0
     V1        0
     V2        0
     V3        0
     V4        0
     V5        0
     V6        0
     V7        0
     V8        0
     V9        0
     V10       0
     V11       0
     V12       0
     V13       0
     V14       0
     V15       0
     V16       0
     V17       0
     V18       0
     V19       0
     V20       0
     V21       0
     V22       0
     V23       0
     V24       0
     V25       0
     V26       0
     V27       0
     V28       0
     Amount    0
```

```
class      0
dtype: int64
```

### #Multicollinearity check with VIF

Multicollinearity occurs when there are two or more independent variables in a multiple regression model, which have a high correlation among themselves. When some features are highly correlated, we might have difficulty in distinguishing between their individual effects on the dependent variable. Multicollinearity can be detected using various techniques, one such technique being the Variance Inflation Factor(VIF).

```
[ ]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[ ]: X = creditcard_df[['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13',
                        'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21']]

vif_data = pd.DataFrame()
vif_data["feature"] = X.columns

vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                   for i in range(len(X.columns))]
```

```
[ ]: print(vif_data)
```

	feature	VIF
0	Time	2.339084
1	V1	1.621694
2	V2	3.869377
3	V3	1.255585
4	V4	1.137944
5	V5	2.753075
6	V6	1.522122
7	V7	2.510165
8	V8	1.097151
9	V9	1.018831
10	V10	1.115668
11	V11	1.028861
12	V12	1.011961
13	V13	1.003434
14	V14	1.026832
15	V15	1.014135
16	V16	1.000371
17	V17	1.004772
18	V18	1.006568
19	V19	1.037809
20	V20	2.233934
21	V21	1.100720

```

22     V22     1.082384
23     V23     1.149268
24     V24     1.000659
25     V25     1.013388
26     V26     1.000487
27     V27     1.008979
28     V28     1.001425
29 Amount    11.499791

```

```

[ ]: def plot_heatmap_of_correlation(df):
    corr = df.corr()
    plt.figure(figsize=(18, 10))
    heat = sns.heatmap(data=corr, annot=True)
    plt.title('Heatmap of Correlation')
    plt.show()

def plot_fraud_transaction_histogram(df):
    fraud = df[df['class'] == 1]
    plt.figure(figsize=(8, 6))
    sns.histplot(data=fraud, x='Amount', bins=65, kde=True)
    plt.title("Distribution of Fraud Transactions")
    plt.xlabel('Amount')
    plt.ylabel('Frequency')
    plt.show()

def plot_real_transaction_histogram(df):
    real = df[df['class'] == 0]
    plt.figure(figsize=(8, 6))
    sns.histplot(data=real, x='Amount', bins=65, kde=True)
    plt.title("Distribution of Real Transactions")
    plt.xlabel('Amount')
    plt.ylabel('Frequency')
    plt.show()

def plot_amount_distribution_for_transactions(df):
    plt.figure(figsize=(6, 4))
    sns.boxplot(x='class', y='Amount', data=df)
    plt.title('Amount Distribution for Fraud and Real Transactions')
    plt.show()

def plot_transaction_time_distribution(df):
    plt.figure(figsize=(6, 4))
    plt.title('Distribution of Transactions with Respect to Time', fontsize=14)
    sns.distplot(df['Time'], bins=120)
    plt.show()

def plot_transaction_counts_by_hour(df):

```

```

plt.figure(figsize=(10, 6))
df['Hour'] = df['Time'] // 3600
plt.title('Transaction Counts by Hour', fontsize=14)
sns.countplot(x='Hour', hue='class', data=df)
plt.xlabel('Hour of the Day')
plt.ylabel('Transaction Count')
plt.xticks(rotation=45)
plt.show()

def plot_transaction_counts_by_day(df):
    plt.figure(figsize=(10, 6))
    df['Day'] = 1
    df.loc[df['Time'] >= 25*3600, 'Day'] = 2
    plt.title('Transaction Counts by Day', fontsize=14)
    sns.countplot(x='Day', hue='class', data=df)
    plt.xlabel('Day')
    plt.ylabel('Transaction Count')
    plt.show()

def plot_transaction_amount_vs_time(df):
    plt.figure(figsize=(10, 6))
    plt.scatter(df[df['class'] == 0]['Time'], df[df['class'] == 0]['Amount'],
        label='Real', alpha=0.5, marker='o')
    plt.scatter(df[df['class'] == 1]['Time'], df[df['class'] == 1]['Amount'],
        label='Fraud', alpha=0.5, marker='x')
    plt.title('Transaction Amount vs Time')
    plt.xlabel('Time')
    plt.ylabel('Amount')
    plt.legend()
    plt.show()


def plot_transaction_amount_vs_features(df):
    features = ['V1', 'V2', 'V3', 'V4', 'V5']
    plt.figure(figsize=(14, 10))
    for i, feature in enumerate(features, start=1):
        plt.subplot(2, 3, i)
        sns.scatterplot(x=feature, y='Amount', hue='class', data=df, alpha=0.5)
        plt.title(f'Transaction Amount vs {feature}')
        plt.xlabel(feature)
        plt.ylabel('Amount')
    plt.tight_layout()
    plt.show()

def plot_class_distribution(df):
    class_counts = df['class'].value_counts()
    plt.figure(figsize=(8, 6))

```

```
plt.pie(class_counts, labels=class_counts.index, autopct='%1.1f%%',
↪startangle=140)
plt.title('Transaction Class Distribution')
plt.axis('equal')
plt.show()
```

```
[ ]: # See distribution of target class
creditcard_df['class'].value_counts().plot.bar(figsize=(4,4),
↪xlabel='Fraudulent', ylabel='Count', title='Credit Card Transactions Data_
↪Distribution')
plt.show()
```



```
[ ]: def dataset_analysis(df):
    fraud = df[df['class'] == 1]
    real = df[df['class'] == 0]

    percent_real = (len(real) / (len(real) + len(fraud))) * 100
    percent_fraud = (len(fraud) / (len(real) + len(fraud))) * 100

    print(f'Number of Real Transactions = {len(real)} and the Percentage of_
↪Real Transactions = {percent_real:.3f}%')
    print(f'Number of Fraud Transactions = {len(fraud)} and the Percentage of_
↪Fraud Transactions = {percent_fraud:.3f}%')
```



```
dataset_analysis(creditcard_df)
```

Number of Real Transactions = 284315 and the Percentage of Real Transactions = 99.827%

Number of Fraud Transactions = 492 and the Percentage of Fraud Transactions = 0.173%

In the below pie chart we are visually representing the distribution of transaction classes (real vs. fraud) in the dataset.


Pie chart Analysis -

1. The pie chart visually shows the proportion of real transactions compared to fraudulent transactions in the dataset.
2. The larger portion of the pie represents real transactions, while the smaller portion represents fraudulent transactions.
3. Each slice of the pie is labeled with the corresponding class (real or fraud) and its percentage of the total number of transactions.

There are 284,315 real transactions, accounting for approximately 99.827% of the total.

There are 492 fraudulent transactions, accounting for approximately 0.173% of the total.

```
[ ]: plot_class_distribution(creditcard_df)
```




The below correlation matrix shows how each feature in the dataset correlates with every other feature.

1. It is useful for understanding the relationships between different features in a dataset.
2. A high positive correlation (close to 1) between two features indicates that they tend to increase or decrease together.
3. A high negative correlation (close to -1) indicates that as one feature increases, the other tends to decrease, and vice versa.
4. A correlation close to 0 suggests that there is little to no linear relationship between the features.
5. By analyzing the heatmap, you can identify patterns and dependencies between features, which can be helpful in feature selection, dimensionality reduction, and understanding the underlying structure of the data.

From below plot, we can conclude that the features are not correlated.


```
[ ]: # Plot heatmap of correlations
plot_heatmap_of_correlation(creditcard_df)
```




Below histograms provide a visual representation of the distribution of transaction amounts for both fraud and real transactions.

1. By comparing the two histograms, you can visually assess whether there are any noticeable differences in the distribution of transaction amounts between fraud and real transactions.
2. Understanding these differences can help in identifying potential patterns or anomalies associated with fraudulent transactions, such as unusually high or low transaction amounts.
3. This analysis can be valuable for developing fraud detection algorithms or understanding the characteristics of fraudulent transactions in the dataset.

```
[ ]: plot_fraud_transaction_histogram(creditcard_df)
```



```
[ ]: plot_real_transaction_histogram(creditcard_df)
```



Above Histogram plot result analysis

```
[ ]: print("Fraud Transaction distribution:")
print(creditcard_df[creditcard_df['class'] == 1]['Amount'].value_counts().
      ↪head())
print("\n")
print("Maximum amount of fraud transaction:",
      ↪creditcard_df[creditcard_df['class'] == 1]['Amount'].max())
print("Minimum amount of fraud transaction:",
      ↪creditcard_df[creditcard_df['class'] == 1]['Amount'].min())
```

Fraud Transaction distribution:

```
1.00    113
0.00     27
99.99    27
0.76     17
0.77     10
```

Name: Amount, dtype: int64


Maximum amount of fraud transaction: 2125.87

Minimum amount of fraud transaction: 0.0

The purpose of the below plot is to visualize and compare the distribution of transaction amounts between fraud and real transactions using a boxplot.

1. This plot helps in identifying any significant differences or patterns in the transaction amounts associated with fraudulent activities compared to legitimate transactions.
2. Moreover, Understanding these differences can aid in developing effective fraud detection strategies and models.

```
[ ]: plot_amount_distribution_for_transactions(creditcard_df)
```



Below plot is to visualize the distribution of transactions over time.

1. The histogram provides insights into the frequency or density of transactions occurring at different points in time.
2. Peaks in the histogram indicate periods of higher transaction activity, while valleys represent periods of lower activity.
3. By examining the shape and pattern of the histogram, you can identify any temporal trends or patterns in transaction activity.

Overall, plot helps in understanding the overall pattern of transaction activity, such as peak hours or periods of increased transaction volume.

```
[ ]: plot_transaction_time_distribution(creditcard_df)
```


<ipython-input-40-ad04cba562f0>:35: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Time'], bins=120)
```




Below plot to visualize the distribution of transaction counts by hour of the day, categorized by transaction class (fraudulent vs. non-fraudulent).

1. The x-axis represents the hours of the day, while the y-axis represents the count of transactions.
2. The bars are color-coded based on the transaction class (fraudulent vs. non-fraudulent), allowing for easy comparison between the two classes.
3. By examining the plot, you can identify any patterns or trends in transaction counts through-

out the day and observe if there are specific hours with higher or lower transaction activity.

Overall, plot helps in understanding the temporal patterns of transaction activity and identifying any anomalies or trends associated with specific hours.

```
[ ]: plot_transaction_counts_by_hour(creditcard_df)
```




Below plot to visualize the distribution of transaction counts by day, categorized by transaction class (fraudulent vs. non-fraudulent)

1. The x-axis represents the days, while the y-axis represents the count of transactions.
2. The bars are color-coded based on the transaction class (fraudulent vs. non-fraudulent), allowing for easy comparison between the two classes.
3. By examining the plot, you can identify any patterns or trends in transaction counts throughout the days and observe if there are specific days with higher or lower transaction activity.

Overall, plot helps in understanding the temporal patterns of transaction activity and identifying any anomalies or trends associated with specific days.

```
[ ]: plot_transaction_counts_by_day(creditcard_df)
```



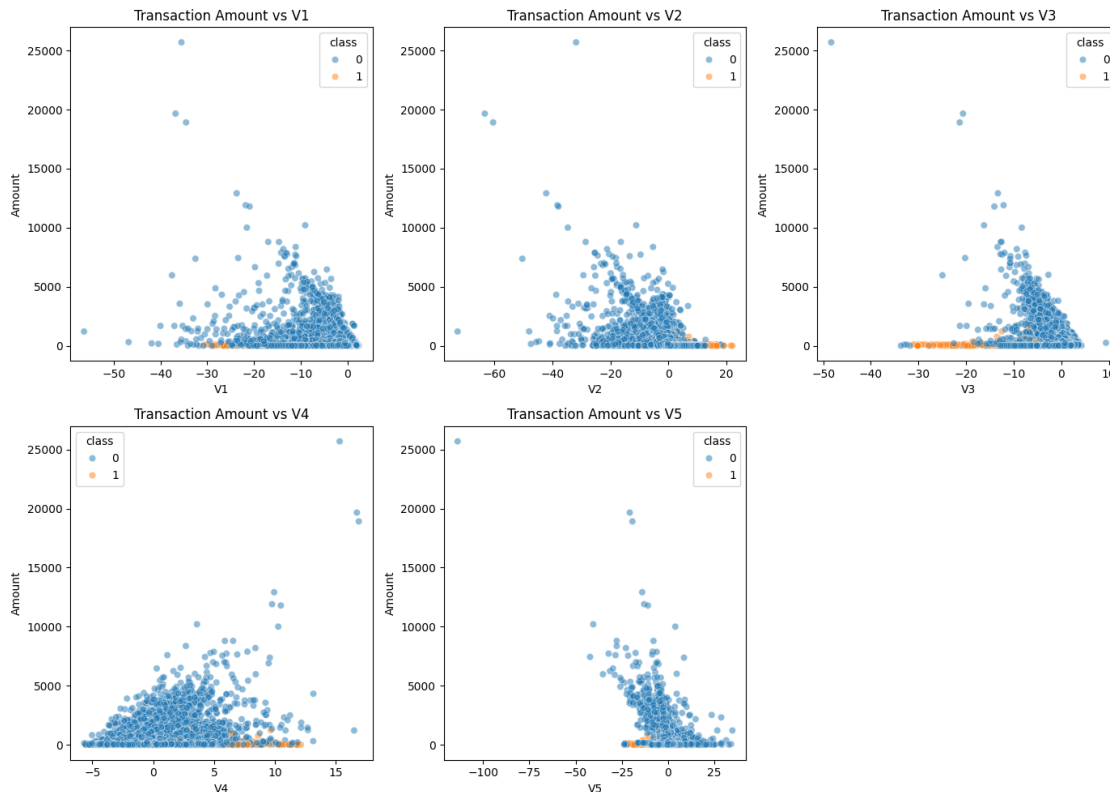
Below code generates multiple scatter plots to visualize the relationship between transaction amounts and selected features (V1 to V5), categorized by transaction class (fraudulent vs. non-fraudulent).

1. Each scatter plot shows the relationship between transaction amounts ('Amount') and a specific feature (V1 to V5).
2. The x-axis represents the feature values, while the y-axis represents the transaction amounts. Points in the scatter plot represent individual transactions.
3. Each subplot is color-coded based on the transaction class (fraudulent vs. non-fraudulent), with different colors representing different classes.
4. By examining the scatter plots, you can observe how transaction amounts vary concerning different feature values for both fraud and non-fraud transactions. The scatter plots allow you to identify any clusters, outliers, or correlations between transaction amounts and the selected features.

Overall, the plot generated provides valuable insights into the relationship between transaction amounts and selected features, aiding in the analysis and detection of fraudulent transactions.

```
[ ]: plot_transaction_amount_vs_features(creditcard_df)
```





## Class Imbalance

Whenever there is an issue with a classification problem and the classes are not represented equally, we say that the data is imbalanced. Applying classifiers to the dataset would most likely result in inaccurate predictions for every category. This was seen as a hurdle when attempting to learn from extremely Imbalanced data sets.

**SMOTE** stands for “Synthetic Minority OverSampling Method.” In this method, synthesis new data from the minority class rather than replicating existing data.

SMOTE creates synthetic samples of the minority class by selecting similar instances and creating new synthetic examples along the line segments joining those instances in the feature space.

```
[ ]: X = creditcard_df.drop('class', axis=1)
      y = creditcard_df['class']

      smote = SMOTE(random_state=42)
```

```
[ ]: X_resampled, y_resampled = smote.fit_resample(X, y)

      balanced_df = X_resampled.copy()
      balanced_df['class'] = y_resampled
```

Balanced dataframe after SMOTE

```
[ ]: balanced_df.head()
```

```
[ ]:
   Time      V1      V2      V3      V4      V5      V6      V7 \
0  0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1  0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2  1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3  1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4  2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941


      V8      V9  ...      V21      V22      V23      V24      V25 \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -0.206010

      V26      V27      V28  Amount  class
0 -0.189115  0.133558 -0.021053   149.62      0
1  0.125895 -0.008983  0.014724     2.69      0
2 -0.139097 -0.055353 -0.059752   378.66      0
3 -0.221929  0.062723  0.061458   123.50      0
4  0.502292  0.219422  0.215153    69.99      0
```

[5 rows x 31 columns]


### #Dataset After Processing

```
[ ]: # See distribution of target class
balanced_df['class'].value_counts().plot.bar(figsize=(4,4),
        xlabel='Fraudulent', ylabel='Count', title='Credit Card Transactions Data_
        Distribution')
plt.show()
```




#### Plots after balancing the datasets


```
[ ]: plot_transaction_counts_by_hour(balanced_df)
```



```
[ ]: plot_transaction_counts_by_day(balanced_df)
```




```
[ ]: plot_transaction_amount_vs_time(balanced_df)
```

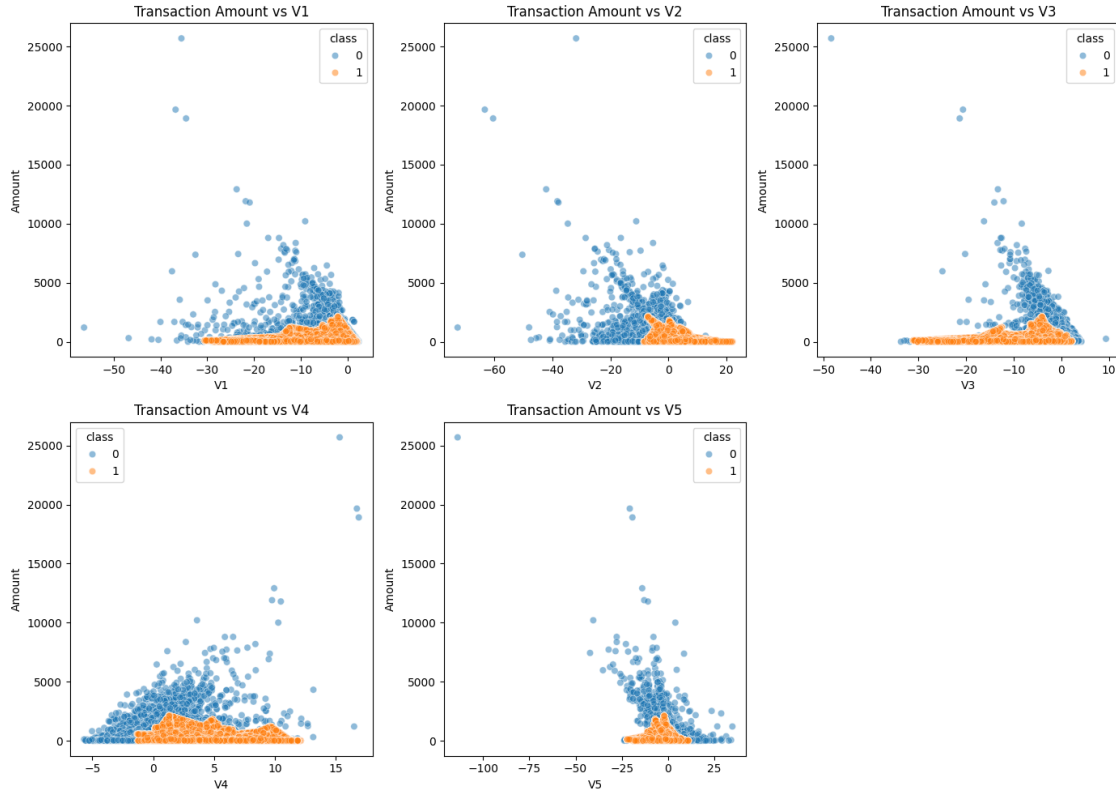


```
[ ]: plot_class_distribution(balanced_df)
```

Transaction Class Distribution



```
[ ]: plot_transaction_amount_vs_features(balanced_df)
```



#Dataset before processing

```
[ ]: creditcard_df
```

```
[ ]:
```

	Time	V1	V2	V3	V4	V5	\	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321		
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018		
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198		
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309		
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193		
...	...	...	...	...	...	...		
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473		
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229		
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515		
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961		
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546		
	V6	V7	V8	V9	...	V21	V22	\
0	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	
1	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	
2	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	
3	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	

```

4      0.095921  0.592941 -0.270533  0.817739 ... -0.009431  0.798278
...
284802 -2.606837 -4.918215  7.305334  1.914428 ...  0.213454  0.111864
284803  1.058415  0.024330  0.294869  0.584800 ...  0.214205  0.924384
284804  3.031260 -0.296827  0.708417  0.432454 ...  0.232045  0.578229
284805  0.623708 -0.686180  0.679145  0.392087 ...  0.265245  0.800049
284806 -0.649617  1.577006 -0.414650  0.486180 ...  0.261057  0.643078

```

```

      V23      V24      V25      V26      V27      V28  Amount  \
0      -0.110474  0.066928  0.128539 -0.189115  0.133558 -0.021053  149.62
1      0.101288 -0.339846  0.167170  0.125895 -0.008983  0.014724   2.69
2      0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752  378.66
3     -0.190321 -1.175575  0.647376 -0.221929  0.062723  0.061458  123.50
4     -0.137458  0.141267 -0.206010  0.502292  0.219422  0.215153   69.99
...
284802  1.014480 -0.509348  1.436807  0.250034  0.943651  0.823731    0.77
284803  0.012463 -1.016226 -0.606624 -0.395255  0.068472 -0.053527   24.79
284804 -0.037501  0.640134  0.265745 -0.087371  0.004455 -0.026561   67.88
284805 -0.163298  0.123205 -0.569159  0.546668  0.108821  0.104533   10.00
284806  0.376777  0.008797 -0.473649 -0.818267 -0.002415  0.013649  217.00

```

```

      class
0         0
1         0
2         0
3         0
4         0
...
284802    0
284803    0
284804    0
284805    0
284806    0

```

[284807 rows x 31 columns]

**#Dataset after processing**

```
[ ]: balanced_df
```

```

[ ]:      Time      V1      V2      V3      V4      V5  \
0      0.000000 -1.359807 -0.072781  2.536347  1.378155 -0.338321
1      0.000000  1.191857  0.266151  0.166480  0.448154  0.060018
2      1.000000 -1.358354 -1.340163  1.773209  0.379780 -0.503198
3      1.000000 -0.966272 -0.185226  1.792993 -0.863291 -0.010309
4      2.000000 -1.158233  0.877737  1.548718  0.403034 -0.407193
...

```



568625	144838.659385	-6.379157	1.672637	-5.885670	2.068340	-0.668576
568626	65965.011763	-2.479028	0.958932	-1.782249	1.541783	-1.191990
568627	34592.129093	-1.799894	2.368957	-2.673997	1.705968	-1.355923
568628	129683.002907	0.255234	2.432041	-5.388252	3.793925	-0.230814
568629	91471.277869	-4.453646	3.210469	-5.294410	1.449911	-1.264653

	V6	V7	V8	V9	...	V21	V22	\
0	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	
1	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	
2	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	
3	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	
4	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	
...	...	...	...	...	...	...	...	
568625	-3.336450	-4.995823	2.632847	-2.275158	...	0.641337	-0.249308	
568626	-0.466794	-1.957161	0.312580	-0.433956	...	0.351983	0.208869	
568627	-1.121788	-2.057832	-1.677459	-0.659287	...	1.473371	-0.581778	
568628	-1.382725	-1.572929	0.748305	-1.600633	...	0.316760	-0.036858	
568629	-0.493626	-3.130644	-4.165957	0.998760	...	4.414468	-1.065864	

	V23	V24	V25	V26	V27	V28	\
0	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	
1	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	
2	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	
3	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	
4	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	
...	...	...	...	...	...	...	
568625	-2.311290	-0.159402	1.190079	-0.258067	0.777265	-0.728919	
568626	-0.235986	-0.404446	0.220454	0.685263	-0.890346	0.598736	
568627	-0.013899	-0.144597	0.120315	0.242272	-0.121166	-0.534238	
568628	0.182968	0.190701	-0.339250	-0.272824	0.315507	-0.091005	
568629	0.798149	0.299668	0.064660	-0.446730	-0.363233	1.018147	

	Amount	class
0	149.620000	0
1	2.690000	0
2	378.660000	0
3	123.500000	0
4	69.990000	0
...	...	...
568625	7.334751	1
568626	74.507571	1
568627	102.486823	1
568628	58.346854	1
568629	143.872749	1

[568630 rows x 31 columns]

## #Model Implementation

Implmented following 7 models on imbalanced data and balanced data achieved with the help of SMOTE algorithm.

1. Naive Bayes
2. Logistic Regression
3. Decision Tree
4. Random Forest
5. Long short-term memory (LSTM)
6. k-nearest neighbors (KNN)
7. Feedforward neural Network (FNN)

## #Imbalanced dataset

```
[ ]: from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, classification_report, precision_score, recall_score, roc_curve, auc
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import LSTM, Dense, Dropout
from keras.optimizers import Adam
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
[ ]: X = creditcard_df.drop(columns=['class'])
y = creditcard_df['class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## #Naive Bayes

```
[ ]: params = {'var_smoothing': np.logspace(0, -9, num=100)}
nb_classifier = GaussianNB()
grid_search = GridSearchCV(nb_classifier, params, cv=5, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

best_nb_classifier = grid_search.best_estimator_

y_pred = best_nb_classifier.predict(X_test)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
/usr/local/lib/python3.10/dist-packages/joblib/externals/loky/backend/fork_exec.py:38: RuntimeWarning:
os.fork() was called. os.fork() is incompatible with multithreaded code, and JAX
```

is multithreaded, so this will likely lead to a deadlock.  
pid = os.fork()

```
[ ]: accuracy_NB = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_NB)

conf_matrix_NB = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_NB, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()


f1_NB = f1_score(y_test, y_pred)
print("F1 Score:", f1_NB)

precision_NB = precision_score(y_test, y_pred)
print("Precision:", precision_NB)

recall_NB = recall_score(y_test, y_pred)
print("Recall:", recall_NB)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9986482216214319



F1 Score: 0.37398373983739835

Precision: 0.92

Recall: 0.23469387755102042

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.92	0.23	0.37	98
accuracy			1.00	56962
macro avg	0.96	0.62	0.69	56962
weighted avg	1.00	1.00	1.00	56962


```
[ ]: y_prob = best_nb_classifier.predict_proba(X_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_NB = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_NB)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
↪roc_auc_NB)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.9453806001860509



### #Logistic Regression

```
[ ]: param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.001, 0.01, 0.1, 1, 10, 100]
}

logreg = LogisticRegression()

grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid, cv=5,
    ↳scoring='accuracy', verbose=1, n_jobs=-1)

grid_search.fit(X_train, y_train)

best_logreg = grid_search.best_estimator_

y_pred = best_logreg.predict(X_test)

print("\nBest Hyperparameters:", grid_search.best_params_)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```
/usr/local/lib/python3.10/dist-  
packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:  
30 fits failed out of a total of 60.  
The score on these train-test partitions for these parameters will be set to  
nan.  
If these failures are not expected, you can try to debug them by setting  
error_score='raise'.
```

Below are more details about the failures:

```
-----  
30 fits failed with the following error:  
Traceback (most recent call last):  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/linear_model/_logistic.py", line 1162, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver  
    raise ValueError(  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.  
  
    warnings.warn(some_fits_failed_message, FitFailedWarning)  
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952:  
UserWarning: One or more of the test scores are non-finite: [      nan  
0.99877987      nan 0.99891593      nan 0.9990476  
      nan 0.99898615      nan 0.99895543      nan 0.99893349]  
    warnings.warn(  

```

Best Hyperparameters: {'C': 0.1, 'penalty': 'l2'}

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html  
Please also refer to the documentation for alternative solver options:  
https://scikit-learn.org/stable/modules/linear_model.html#logistic-  
regression  
n_iter_i = _check_optimize_result(  

```

```
[ ]: accuracy_LR = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy_LR)
```

```

conf_matrix_LR = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_LR, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_LR = f1_score(y_test, y_pred)
print("F1 Score:", f1_LR)


precision_LR = precision_score(y_test, y_pred)
print("Precision:", precision_LR)

recall_LR = recall_score(y_test, y_pred)
print("Recall:", recall_LR)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9986833327481479



F1 Score: 0.6113989637305699

Precision: 0.6210526315789474

Recall: 0.6020408163265306

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.62	0.60	0.61	98

accuracy			1.00	56962
macro avg	0.81	0.80	0.81	56962
weighted avg	1.00	1.00	1.00	56962

```
[ ]: y_prob = best_logreg.predict_proba(X_test)[:, 1]


fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_LR = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_LR)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %_
↪roc_auc_LR)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.9050803636029539





### #Decision Tree

```
[ ]: clf = DecisionTreeClassifier(max_depth=15, max_features='auto',
    ↳ criterion='entropy', random_state=42)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set
`max_features='sqrt'`.
  warnings.warn(
```

```
[ ]: accuracy_DT = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_DT)

conf_matrix_DT = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
```

```

sns.heatmap(conf_matrix_DT, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_DT = f1_score(y_test, y_pred)
print("F1 Score:", f1_DT)

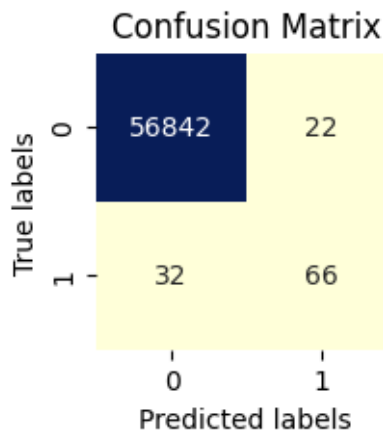
precision_DT = precision_score(y_test, y_pred)
print("Precision:", precision_DT)

recall_DT = recall_score(y_test, y_pred)
print("Recall:", recall_DT)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9990519995786665



F1 Score: 0.7096774193548386

Precision: 0.75

Recall: 0.673469387755102

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.75	0.67	0.71	98
accuracy			1.00	56962
macro avg	0.87	0.84	0.85	56962

weighted avg      1.00      1.00      1.00      56962


```
[ ]: y_prob = clf.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_DT = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_DT)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
        ↪roc_auc_DT)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.8342387637384723



## #Random Forest

```
[ ]: clf = RandomForestClassifier(n_estimators=50, random_state=42)

clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

[ ]: accuracy_RF = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_RF)

conf_matrix_RF = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_RF, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()


f1_RF = f1_score(y_test, y_pred)
print("F1 Score:", f1_RF)

precision_RF = precision_score(y_test, y_pred)
print("Precision:", precision_RF)

recall_RF = recall_score(y_test, y_pred)
print("Recall:", recall_RF)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9995435553526912



F1 Score: 0.8505747126436782  
Precision: 0.9736842105263158  
Recall: 0.7551020408163265

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.97	0.76	0.85	98
accuracy			1.00	56962
macro avg	0.99	0.88	0.93	56962
weighted avg	1.00	1.00	1.00	56962


```
[ ]: y_prob = clf.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_RF = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_RF)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc_RF)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.9480902518576366



### #Long Short-Term Memory Networks (LSTM)

```
[ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

timesteps = 1
X_train_reshaped = X_train_scaled.reshape(X_train_scaled.shape[0], timesteps,
    ↳X_train_scaled.shape[1])
X_test_reshaped = X_test_scaled.reshape(X_test_scaled.shape[0], timesteps,
    ↳X_test_scaled.shape[1])

model = Sequential()
model.add(LSTM(units=64, input_shape=(X_train_reshaped.shape[1],
    ↳X_train_reshaped.shape[2])))
model.add(Dropout(0.2))
model.add(Dense(units=1, activation='sigmoid'))

optimizer = Adam(lr=0.001)
```

```

model.compile(optimizer=optimizer, loss='binary_crossentropy',
↳metrics=['accuracy'])

model.fit(X_train_resaped, y_train, epochs=10, batch_size=64,
↳validation_split=0.2)

y_pred_proba = model.predict(X_test_resaped)
y_pred = (y_pred_proba > 0.5).astype(int)

```

WARNING:absl:lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.Adam.

```

Epoch 1/10
2849/2849 [=====] - 15s 4ms/step - loss: 0.0317 -
accuracy: 0.9966 - val_loss: 0.0028 - val_accuracy: 0.9994
Epoch 2/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0030 -
accuracy: 0.9994 - val_loss: 0.0025 - val_accuracy: 0.9994
Epoch 3/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0027 -
accuracy: 0.9995 - val_loss: 0.0024 - val_accuracy: 0.9995
Epoch 4/10
2849/2849 [=====] - 11s 4ms/step - loss: 0.0024 -
accuracy: 0.9995 - val_loss: 0.0024 - val_accuracy: 0.9994
Epoch 5/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0022 -
accuracy: 0.9995 - val_loss: 0.0026 - val_accuracy: 0.9994
Epoch 6/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0021 -
accuracy: 0.9995 - val_loss: 0.0026 - val_accuracy: 0.9995
Epoch 7/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0020 -
accuracy: 0.9995 - val_loss: 0.0025 - val_accuracy: 0.9994
Epoch 8/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0019 -
accuracy: 0.9996 - val_loss: 0.0024 - val_accuracy: 0.9994
Epoch 9/10
2849/2849 [=====] - 13s 4ms/step - loss: 0.0018 -
accuracy: 0.9996 - val_loss: 0.0025 - val_accuracy: 0.9994
Epoch 10/10
2849/2849 [=====] - 12s 4ms/step - loss: 0.0016 -
accuracy: 0.9996 - val_loss: 0.0027 - val_accuracy: 0.9995
1781/1781 [=====] - 4s 2ms/step

```

```

[ ]: accuracy_LSTM = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_LSTM)

```

```

conf_matrix_LSTM = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_LSTM, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_LSTM = f1_score(y_test, y_pred)
print("F1 Score:", f1_LSTM)


precision_LSTM = precision_score(y_test, y_pred)
print("Precision:", precision_LSTM)

recall_LSTM = recall_score(y_test, y_pred)
print("Recall:", recall_LSTM)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.999403110845827



F1 Score: 0.8068181818181819

Precision: 0.9102564102564102

Recall: 0.7244897959183674

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.91	0.72	0.81	98




accuracy			1.00	56962
macro avg	0.95	0.86	0.90	56962
weighted avg	1.00	1.00	1.00	56962

```
[ ]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculating AUC-ROC
roc_auc_LSTM = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_LSTM)

# Plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %_
↪roc_auc_LSTM)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.9826097785765967



#KNN

```
[ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9995259997893332

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.94	0.78	0.85	98
accuracy			1.00	56962
macro avg	0.97	0.89	0.92	56962
weighted avg	1.00	1.00	1.00	56962

```
[ ]: accuracy_KNN = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_KNN)

conf_matrix_KNN = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_KNN, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()


f1_KNN = f1_score(y_test, y_pred)
print("F1 Score:", f1_KNN)

precision_KNN = precision_score(y_test, y_pred)
print("Precision:", precision_KNN)

recall_KNN = recall_score(y_test, y_pred)
print("Recall:", recall_KNN)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9995259997893332



F1 Score: 0.8491620111731844  
Precision: 0.9382716049382716  
Recall: 0.7755102040816326

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.94	0.78	0.85	98
accuracy			1.00	56962
macro avg	0.97	0.89	0.92	56962
weighted avg	1.00	1.00	1.00	56962

```
[ ]: y_pred_proba = knn.predict_proba(X_test_scaled)[: , 1]


fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

roc_auc_KNN = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_KNN)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
    ↪roc_auc_KNN)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
```

```
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.9335599690776705



#FNN

```
[ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Defining the FNN model
model = Sequential()
model.add(Dense(units=64, activation='relu', input_shape=(X_train_scaled.
    ↪shape[1],)))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=1, activation='sigmoid'))

# Compiling the model
```

```

model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy'])

# Training the model
model.fit(X_train_scaled, y_train, epochs=10, batch_size=64, validation_split=0.
↳2)

# Evaluating the model
y_pred_proba = model.predict(X_test_scaled)
y_pred = (y_pred_proba > 0.5).astype(int)

```

```

Epoch 1/10
2849/2849 [=====] - 9s 3ms/step - loss: 0.0110 -
accuracy: 0.9976 - val_loss: 0.0032 - val_accuracy: 0.9994
Epoch 2/10
2849/2849 [=====] - 8s 3ms/step - loss: 0.0033 -
accuracy: 0.9994 - val_loss: 0.0030 - val_accuracy: 0.9993
Epoch 3/10
2849/2849 [=====] - 6s 2ms/step - loss: 0.0029 -
accuracy: 0.9994 - val_loss: 0.0026 - val_accuracy: 0.9994
Epoch 4/10
2849/2849 [=====] - 8s 3ms/step - loss: 0.0025 -
accuracy: 0.9995 - val_loss: 0.0027 - val_accuracy: 0.9993
Epoch 5/10
2849/2849 [=====] - 6s 2ms/step - loss: 0.0023 -
accuracy: 0.9995 - val_loss: 0.0029 - val_accuracy: 0.9993
Epoch 6/10
2849/2849 [=====] - 7s 3ms/step - loss: 0.0021 -
accuracy: 0.9995 - val_loss: 0.0031 - val_accuracy: 0.9994
Epoch 7/10
2849/2849 [=====] - 7s 2ms/step - loss: 0.0021 -
accuracy: 0.9995 - val_loss: 0.0026 - val_accuracy: 0.9995
Epoch 8/10
2849/2849 [=====] - 8s 3ms/step - loss: 0.0018 -
accuracy: 0.9996 - val_loss: 0.0026 - val_accuracy: 0.9995
Epoch 9/10
2849/2849 [=====] - 7s 2ms/step - loss: 0.0018 -
accuracy: 0.9996 - val_loss: 0.0033 - val_accuracy: 0.9993
Epoch 10/10
2849/2849 [=====] - 8s 3ms/step - loss: 0.0016 -
accuracy: 0.9996 - val_loss: 0.0029 - val_accuracy: 0.9996
1781/1781 [=====] - 3s 2ms/step

```

```

[ ]: accuracy_FNN = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_FNN)

conf_matrix_FNN = confusion_matrix(y_test, y_pred)

```

```

plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_FNN, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_FNN = f1_score(y_test, y_pred)
print("F1 Score:", f1_FNN)


precision_FNN = precision_score(y_test, y_pred)
print("Precision:", precision_FNN)

recall_FNN = recall_score(y_test, y_pred)
print("Recall:", recall_FNN)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9994908886626171



F1 Score: 0.8361581920903955

Precision: 0.9367088607594937

Recall: 0.7551020408163265

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.94	0.76	0.84	98
accuracy			1.00	56962

macro avg	0.97	0.88	0.92	56962
weighted avg	1.00	1.00	1.00	56962


```
[ ]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculating AUC-ROC
roc_auc_FNN = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_FNN)

# Plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %_
↪roc_auc_FNN)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.9747706306777072



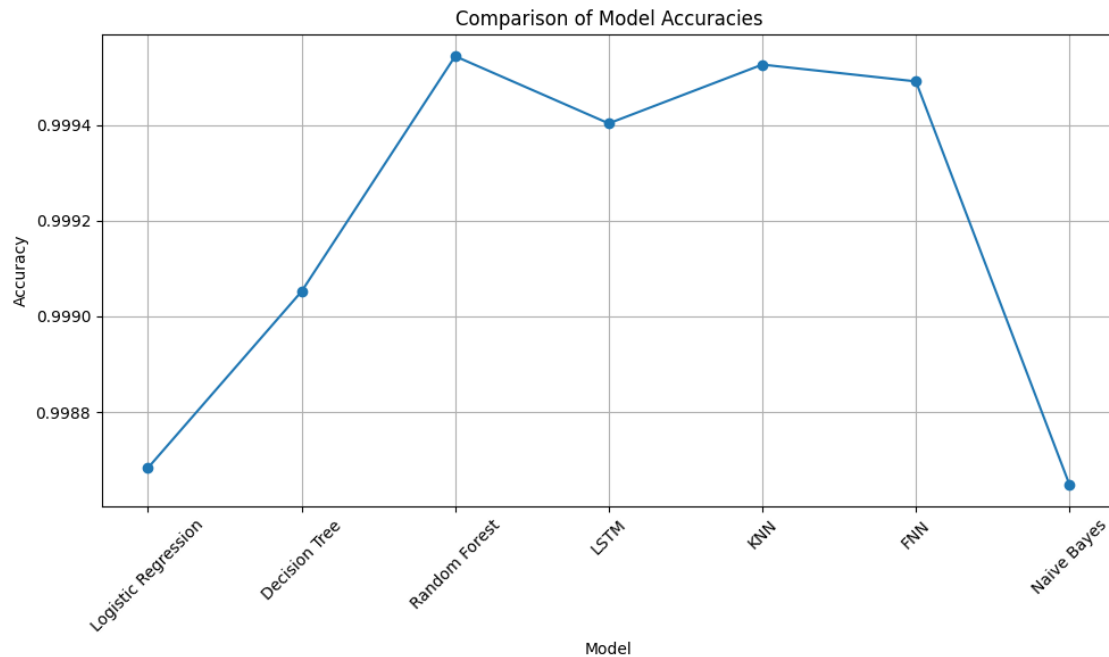


#Comparing Accuracy for all implemented models on Imbalanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM',
          'KNN', 'FNN', 'Naive Bayes']
accuracies = [accuracy_LR, accuracy_DT, accuracy_RF, accuracy_LSTM,
              accuracy_KNN, accuracy_FNN, accuracy_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Accuracies')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



#Comparing F1 - Score for all implemented models on Imbalanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [f1_LR, f1_DT, f1_RF, f1_LSTM, f1_KNN, f1_FNN, f1_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model F1-Scores')
plt.xlabel('Model')
plt.ylabel('F1-Score')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

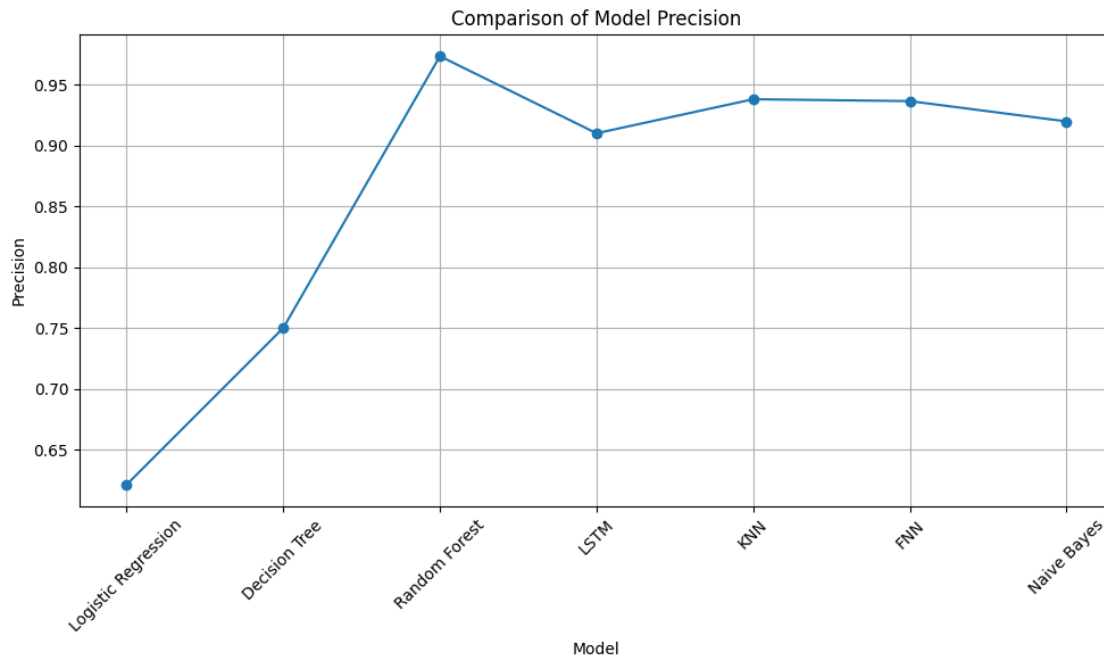


#Comparing Precision for all implemented models on Imbalanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [precision_LR, precision_DT, precision_RF, precision_LSTM, precision_KNN, precision_FNN, precision_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Precision')
plt.xlabel('Model')
plt.ylabel('Precision')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```




#Comparing Recall for all implemented models on Imbalanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [recall_LR, recall_DT, recall_RF, recall_LSTM, recall_KNN, recall_FNN, recall_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Recall')
plt.xlabel('Model')
plt.ylabel('Recall')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```




#Comparing Area under Receiver Operating Characteristic (ROC) Curve for all implemented models on Imbalanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM',
          'KNN', 'FNN', 'Naive Bayes']
accuracies = [roc_auc_LR, roc_auc_DT, roc_auc_RF, roc_auc_LSTM, roc_auc_KNN,
              roc_auc_FNN, roc_auc_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model ROC_AUC')
plt.xlabel('Model')
plt.ylabel('ROC_AUC')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

### #Balanced Dataset

```
[ ]: X = balanced_df.drop(columns=['class'])
y = balanced_df['class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

### #Naive Bayes

```
[ ]: params = {'var_smoothing': np.logspace(0, -9, num=100)}
nb_classifier = GaussianNB()
grid_search = GridSearchCV(nb_classifier, params, cv=5, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

best_nb_classifier = grid_search.best_estimator_
```


```
y_pred = best_nb_classifier.predict(X_test)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
/usr/local/lib/python3.10/dist-  
packages/joblib/externals/loky/backend/fork_exec.py:38: RuntimeWarning:  
os.fork() was called. os.fork() is incompatible with multithreaded code, and JAX  
is multithreaded, so this will likely lead to a deadlock.  
    pid = os.fork()
```

```
[ ]: accuracy_NB = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy_NB)  
  
conf_matrix_NB = confusion_matrix(y_test, y_pred)  
plt.figure(figsize=(2, 2))  
sns.heatmap(conf_matrix_NB, annot=True, cmap="YlGnBu", fmt='g', cbar=False)  
plt.xlabel('Predicted labels')  
plt.ylabel('True labels')  
plt.title('Confusion Matrix')  
plt.show()  
  
f1_NB = f1_score(y_test, y_pred)  
print("F1 Score:", f1_NB)  
  
precision_NB = precision_score(y_test, y_pred)  
print("Precision:", precision_NB)  
  
recall_NB = recall_score(y_test, y_pred)  
print("Recall:", recall_NB)  
  
print("\nClassification Report:")  
print(classification_report(y_test, y_pred))
```

Accuracy: 0.8686404164395125



F1 Score: 0.8504888960057647  
Precision: 0.9894511328970962  
Recall: 0.745752597584948

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.99	0.88	56750
1	0.99	0.75	0.85	56976
accuracy			0.87	113726
macro avg	0.89	0.87	0.87	113726
weighted avg	0.89	0.87	0.87	113726

```
[ ]: y_prob = best_nb_classifier.predict_proba(X_test)[: , 1]


fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_NB = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_NB)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %_
↪roc_auc_NB)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.982717349727283





### #Logistic Regression

```
[ ]: param_grid = {  
    'penalty': ['l1', 'l2'],  
    'C': [0.001, 0.01, 0.1, 1, 10, 100]  
}  
  
logreg = LogisticRegression()  
  
grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid, cv=5,  
    ↪scoring='accuracy', verbose=1, n_jobs=-1)  
  
grid_search.fit(X_train, y_train)  
  
best_logreg = grid_search.best_estimator_  
  
y_pred = best_logreg.predict(X_test)  
  
print("\nBest Hyperparameters:", grid_search.best_params_)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```
/usr/local/lib/python3.10/dist-  
packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:  
30 fits failed out of a total of 60.  
The score on these train-test partitions for these parameters will be set to  
nan.  
If these failures are not expected, you can try to debug them by setting  
error_score='raise'.
```

Below are more details about the failures:

```
-----  
30 fits failed with the following error:  
Traceback (most recent call last):  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/linear_model/_logistic.py", line 1162, in fit  
    solver = _check_solver(self.solver, self.penalty, self.dual)  
  File "/usr/local/lib/python3.10/dist-  
packages/sklearn/linear_model/_logistic.py", line 54, in _check_solver  
    raise ValueError(  
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.  
  
    warnings.warn(some_fits_failed_message, FitFailedWarning)  
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:952:  
UserWarning: One or more of the test scores are non-finite: [      nan  
0.97104884      nan 0.97263818      nan 0.97235461  
      nan 0.97026626      nan 0.9697101      nan 0.96974747]  
    warnings.warn(  

```

Best Hyperparameters: {'C': 0.01, 'penalty': 'l2'}

```
[ ]: accuracy_LR = accuracy_score(y_test, y_pred)  
print("Accuracy:", accuracy_LR)  
  
conf_matrix_LR = confusion_matrix(y_test, y_pred)  
plt.figure(figsize=(2, 2))  
sns.heatmap(conf_matrix_LR, annot=True, cmap="YlGnBu", fmt='g', cbar=False)  
plt.xlabel('Predicted labels')  
plt.ylabel('True labels')  
plt.title('Confusion Matrix')  
plt.show()  
  
f1_LR = f1_score(y_test, y_pred)  
print("F1 Score:", f1_LR)
```

```


precision_LR = precision_score(y_test, y_pred)
print("Precision:", precision_LR)

recall_LR = recall_score(y_test, y_pred)
print("Recall:", recall_LR)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9721259870214375



F1 Score: 0.9719726976941576  
 Precision: 0.9793151368301026  
 Recall: 0.9647395394552092

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.98	0.97	56750
1	0.98	0.96	0.97	56976
accuracy			0.97	113726
macro avg	0.97	0.97	0.97	113726
weighted avg	0.97	0.97	0.97	113726

```

[ ]: y_prob = best_logreg.predict_proba(X_test)[:, 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_LR = auc(fpr, tpr)

```


```

print("ROC_AUC:", roc_auc_LR)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
        ↪roc_auc_LR)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```

ROC\_AUC: 0.9934243790104991



## #Decision Tree

```

[ ]: clf = DecisionTreeClassifier(max_depth=15, max_features='auto',
        ↪criterion='entropy', random_state=42)

```

```
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/tree/_classes.py:269:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set
`max_features='sqrt'`.
  warnings.warn(
```

```
[ ]: accuracy_DT = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_DT)

conf_matrix_DT = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_DT, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()


f1_DT = f1_score(y_test, y_pred)
print("F1 Score:", f1_DT)

precision_DT = precision_score(y_test, y_pred)
print("Precision:", precision_DT)

recall_DT = recall_score(y_test, y_pred)
print("Recall:", recall_DT)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9961662240824438



F1 Score: 0.9961793932596085

Precision: 0.9947324209863148

Recall: 0.9976305812973884

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	56750
1	0.99	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

```
[ ]: y_prob = clf.predict_proba(X_test)[:, 1]


fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_DT = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_DT)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
    ↪roc_auc_DT)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
```

```
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.9988852839498382



### #Random Forest

```
[ ]: clf = RandomForestClassifier(n_estimators=50, random_state=42)
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
[ ]: accuracy_RF = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy_RF)
```

```
conf_matrix_RF = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(2, 2))
```

```
sns.heatmap(conf_matrix_RF, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
```

```
plt.xlabel('Predicted labels')
```

```

plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_RF = f1_score(y_test, y_pred)
print("F1 Score:", f1_RF)


precision_RF = precision_score(y_test, y_pred)
print("Precision:", precision_RF)

recall_RF = recall_score(y_test, y_pred)
print("Recall:", recall_RF)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.999903276295658



F1 Score: 0.9999034774444338

Precision: 0.9998069735202766

Recall: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726




```
[ ]: y_prob = clf.predict_proba(X_test)[: , 1]

fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc_RF = auc(fpr, tpr)

print("ROC_AUC:", roc_auc_RF)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc_RF)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

ROC\_AUC: 0.9999902006811431



## #Long Short-Term Memory Networks (LSTM)

```
[ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

timesteps = 1
X_train_reshaped = X_train_scaled.reshape(X_train_scaled.shape[0], timesteps,
↳X_train_scaled.shape[1])
X_test_reshaped = X_test_scaled.reshape(X_test_scaled.shape[0], timesteps,
↳X_test_scaled.shape[1])

model = Sequential()
model.add(LSTM(units=64, input_shape=(X_train_reshaped.shape[1],
↳X_train_reshaped.shape[2])))
model.add(Dropout(0.2))
model.add(Dense(units=1, activation='sigmoid'))

optimizer = Adam(lr=0.001)
model.compile(optimizer=optimizer, loss='binary_crossentropy',
↳metrics=['accuracy'])

model.fit(X_train_reshaped, y_train, epochs=10, batch_size=64,
↳validation_split=0.2)

y_pred_proba = model.predict(X_test_reshaped)
y_pred = (y_pred_proba > 0.5).astype(int)
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning\_rate` or use the legacy optimizer, e.g.,`tf.keras.optimizers.legacy.Adam`.

```
Epoch 1/10
5687/5687 [=====] - 36s 6ms/step - loss: 0.0614 -
accuracy: 0.9778 - val_loss: 0.0285 - val_accuracy: 0.9887
Epoch 2/10
5687/5687 [=====] - 37s 6ms/step - loss: 0.0237 -
accuracy: 0.9917 - val_loss: 0.0154 - val_accuracy: 0.9947
Epoch 3/10
5687/5687 [=====] - 24s 4ms/step - loss: 0.0157 -
accuracy: 0.9948 - val_loss: 0.0102 - val_accuracy: 0.9968
Epoch 4/10
5687/5687 [=====] - 22s 4ms/step - loss: 0.0116 -
accuracy: 0.9961 - val_loss: 0.0070 - val_accuracy: 0.9983
Epoch 5/10
5687/5687 [=====] - 29s 5ms/step - loss: 0.0089 -
accuracy: 0.9972 - val_loss: 0.0054 - val_accuracy: 0.9986
Epoch 6/10
```

```

5687/5687 [=====] - 24s 4ms/step - loss: 0.0073 -
accuracy: 0.9978 - val_loss: 0.0040 - val_accuracy: 0.9990
Epoch 7/10
5687/5687 [=====] - 26s 5ms/step - loss: 0.0060 -
accuracy: 0.9981 - val_loss: 0.0033 - val_accuracy: 0.9993
Epoch 8/10
5687/5687 [=====] - 32s 6ms/step - loss: 0.0052 -
accuracy: 0.9984 - val_loss: 0.0029 - val_accuracy: 0.9992
Epoch 9/10
5687/5687 [=====] - 25s 4ms/step - loss: 0.0045 -
accuracy: 0.9986 - val_loss: 0.0023 - val_accuracy: 0.9995
Epoch 10/10
5687/5687 [=====] - 22s 4ms/step - loss: 0.0040 -
accuracy: 0.9988 - val_loss: 0.0021 - val_accuracy: 0.9995
3554/3554 [=====] - 7s 2ms/step

```

```

[ ]: accuracy_LSTM = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_LSTM)

conf_matrix_LSTM = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_LSTM, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_LSTM = f1_score(y_test, y_pred)
print("F1 Score:", f1_LSTM)


precision_LSTM = precision_score(y_test, y_pred)
print("Precision:", precision_LSTM)

recall_LSTM = recall_score(y_test, y_pred)
print("Recall:", recall_LSTM)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9994108647099168



F1 Score: 0.999412378638648  
Precision: 0.9988254474694529  
Recall: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726


```
[ ]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculating AUC-ROC
roc_auc_LSTM = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_LSTM)

# Plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
    ↪roc_auc_LSTM)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
```

```
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.999976808536433



#KNN

```
[ ]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

knn = KNeighborsClassifier(n_neighbors=5)

knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9990943144047975

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

```
[ ]: accuracy_KNN = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_KNN)

conf_matrix_KNN = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_KNN, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()


f1_KNN = f1_score(y_test, y_pred)
print("F1 Score:", f1_KNN)

precision_KNN = precision_score(y_test, y_pred)
print("Precision:", precision_KNN)

recall_KNN = recall_score(y_test, y_pred)
print("Recall:", recall_KNN)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9990943144047975



F1 Score: 0.9990969269212222  
Precision: 0.9981954834527584  
Recall: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

```
[ ]: y_pred_proba = knn.predict_proba(X_test_scaled)[: , 1]


fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

roc_auc_KNN = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_KNN)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
    ↪roc_auc_KNN)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
```

```
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.9997092511013216



#FNN

```
[ ]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)


# Calculating AUC-ROC
roc_auc_FNN = auc(fpr, tpr)
print("AUC-ROC:", roc_auc_FNN)

# Plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc_FNN)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```



```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```

AUC-ROC: 0.9997092511013216



```
[ ]: accuracy_FNN = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy_FNN)

conf_matrix_FNN = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(2, 2))
sns.heatmap(conf_matrix_FNN, annot=True, cmap="YlGnBu", fmt='g', cbar=False)
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

f1_FNN = f1_score(y_test, y_pred)
```

```

print("F1 Score:", f1_FNN)


precision_FNN = precision_score(y_test, y_pred)
print("Precision:", precision_FNN)

recall_FNN = recall_score(y_test, y_pred)
print("Recall:", recall_FNN)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.9990943144047975



F1 Score: 0.9990969269212222  
 Precision: 0.9981954834527584  
 Recall: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56750
1	1.00	1.00	1.00	56976
accuracy			1.00	113726
macro avg	1.00	1.00	1.00	113726
weighted avg	1.00	1.00	1.00	113726

```

[ ]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

# Calculating AUC-ROC
roc_auc_FNN = auc(fpr, tpr)

```


```

print("AUC-ROC:", roc_auc_FNN)

# Plotting ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label='ROC curve (AUC = %0.2f)' %
        roc_auc_FNN)
plt.plot([0, 1], [0, 1], color='red', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```

AUC-ROC: 0.9997092511013216




#Comparing Accuracy for all implemented models on balanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [accuracy_LR, accuracy_DT, accuracy_RF, accuracy_LSTM, accuracy_KNN, accuracy_FNN, accuracy_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Accuracies')
plt.xlabel('Model')
plt.ylabel('Accuracy')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```




#Comparing F1 - Score for all implemented models on balanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [f1_LR, f1_DT, f1_RF, f1_LSTM, f1_KNN, f1_FNN, f1_NB]
```

```
plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model F1-Scores')
plt.xlabel('Model')
plt.ylabel('F1-Score')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```




**#Comparing Precision for all implemented models on balanced dataset**

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM', 'KNN', 'FNN', 'Naive Bayes']
accuracies = [precision_LR, precision_DT, precision_RF, precision_LSTM, precision_KNN, precision_FNN, precision_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Precision')
plt.xlabel('Model')
plt.ylabel('Precision')
plt.xticks(rotation=45)
```

```
plt.grid(True)
plt.tight_layout()
plt.show()
```




#Comparing Recall for all implemented models on balanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM',
          'KNN', 'FNN', 'Naive Bayes']
accuracies = [recall_LR, recall_DT, recall_RF, recall_LSTM, recall_KNN,
              recall_FNN, recall_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model Recall')
plt.xlabel('Model')
plt.ylabel('Recall')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



#Comparing Area under Receiver Operating Characteristic (ROC) Curve for all implemented models on balanced dataset

```
[ ]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'LSTM',
          ↪ 'KNN', 'FNN', 'Naive Bayes']
accuracies = [roc_auc_LR, roc_auc_DT, roc_auc_RF, roc_auc_LSTM, roc_auc_KNN,
              ↪ roc_auc_FNN, roc_auc_NB]

plt.figure(figsize=(10, 6))
plt.plot(models, accuracies, marker='o', linestyle='-')
plt.title('Comparison of Model ROC_AUC')
plt.xlabel('Model')
plt.ylabel('ROC_AUC')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

